

Assignment #2 – mdadm Linear Device (Basic Functionality)

CMPSC311 - Introduction to Systems Programming

Spring 2023 - Prof. Aghayev

Due date: February 24, 2023 (11:59 PM) EST

You will need to install `libssl-dev` by running `sudo apt install libssl-dev`.

Today is the first day of your summer internship at a cryptocurrency startup. Before you join, the marketing team decided that they want to differentiate their product by emphasizing on security. On the same day that you join the company, the shipment of 16 military-grade, nuclear bomb-proof hard disks arrives. They are supposed to replace the existing commercial-grade hard disks and will be used to store the most critical user data—cryptocurrency wallets. However, the disk company focuses on physical security and doesn't invest much in software. They provide their disks as a JBOD (Just a Bunch of Disks), which is a storage architecture consisting of numerous disks inside of a single storage enclosure. They also provide a user manual along with the shipment:

Bits	Width	Field	Description
26-31	6	Command	This is the command to be executed by JBOD.
22-25	4	DiskID	This is the ID of the disk to perform operation on
8-21	14	Reserved	Unused bits (for now)
0-7	8	BlockID	Block address within the disk

Table 1: JBOD operation format

Thank you for purchasing our military-grade, nuclear bomb-proof hard disks, built with patented technologies. Each of the disks in front of you consists of 256 blocks, and each block has 256 bytes, coming to a total of $256 \times 256 = 65,536$ bytes per disk. Since you bought 16 disks, the combined capacity is $16 \times 65,536 = 1,048,576$ bytes = 1 MB. We provide you with a device driver with a single function that you can use to control the disks.

```
int jbod_operation(uint32_t op, uint8_t *block);
```

This function returns 0 on success and -1 on failure. It accepts an operation through the `op` parameter, the format of which is described in Table 1, and a pointer to a buffer. The command field can be one of the following commands, which are declared as a C `enum` type in the header that we have provide to you:

1. `JBOD_MOUNT`: mount all disks in the JBOD and make them ready to serve commands. This is the first command that should be called on the JBOD before issuing any other commands; all commands before it will fail. When the command field of `op` is set to this command, all other fields in `op` are ignored by the JBOD driver. Similarly, the `block` argument passed to `jbod_operation` can be `NULL`.
2. `JBOD_UNMOUNT`: unmount all disks in the JBOD. This is the last command that should be called on the JBOD; all commands after it will fail. When the command field of `op` is set to this command, all other fields in `op` are ignored by the JBOD driver. Similarly, the `block` argument passed to `jbod_operation` can be `NULL`.
3. `JBOD_SEEK_TO_DISK`: seeks to a specific disk. JBOD internally maintains an *I/O position*, a tuple consisting of `{CurrentDiskID, CurrentBlockID}`, which determines where the next I/O operation will happen. This command seeks to the beginning of disk specified `DiskID` field in `op`. In other words, it modifies I/O position: it sets `CurrentDiskID` to `DiskID` specified in `op` and it sets `CurrentBlockID` to 0.

When the command field of `op` is set to this command, the `BlockID` field in `op` is ignored by the JBOD driver. Similarly, the `block` argument passed to `jbod_operation` can be `NULL`.

4. `JBOD_SEEK_TO_BLOCK`: seeks to a specific block in current disk. This command sets the `CurrentBlockID` in *I/O position* to the block specified in `BlockID` field in `op`. When the command field of `op` is set to this command, the `DiskID` field in `op` is ignored by the JBOD driver. Similarly, the `block` argument passed to `jbod_operation` can be `NULL`.
5. `JBOD_READ_BLOCK`: reads the block in current I/O position into the buffer specified by the `block` argument to `jbod_operation`. The buffer pointed by `block` must be of block size, that is 256 bytes. **More importantly, after this operation completes, the `CurrentBlockID` in I/O position is incremented by 1; that is, the next I/O operation will happen on the next block of the current disk.** When the command field of `op` is set to this command, all other fields in `op` are ignored by the JBOD driver.
6. `JBOD_WRITE_BLOCK`: writes the data in the `block` buffer into the block in the current I/O position. The buffer pointed by `block` must be of block size, that is 256 bytes. **More importantly, after this operation completes, the `CurrentBlockID` in I/O position is incremented by 1; that is, the next I/O operation will happen on the next block of the current disk.** When the command field of `op` is set to this command, all other fields in `op` are ignored by the JBOD driver.

After you finished your onboarding session with HR and enjoyed the free lunch with your new colleagues, you received an email from the manager of the team.

Welcome, to the team! Here's your task for the next two weeks. You will be working on integrating JBOD into our existing storage system. Specifically, you will implement one of the functionalities of the `mdadm` utility in Linux. `Mdadm` stands for multiple disk and device administration, and it is a tool for doing cool tricks with multiple disks. You will implement one of such tricks supported by `mdadm`, called *linear device*. A linear device makes multiple disks appear as a one large disk to the operating system. In our case, we will use your program to configure 16 disks of size 64 KB as a single 1 MB disk. Below are the functions you need to implement.

`int mdadm_mount(void)`: Mount the linear device; now `mdadm` user can run read and operations on the linear address space that combines all disks. It should return 1 on success and -1 on failure. Calling this function the second time without calling `mdadm_unmount` in between, should fail.

`int mdadm_unmount(void)`: Unmount the linear device; now all commands to the linear device should fail. It should return 1 on success and -1 on failure. Calling this function the second time without calling `mdadm_mount` in between, should fail.

`int mdadm_read(uint32_t addr, uint32_t len, uint8_t *buf)`: Read `len` bytes into `buf` starting at `addr`. Read from an out-of-bound linear address should fail. A read larger than 1,024 bytes should fail; in other words, `len` can be 1,024 at most. There are a few more restrictions that you will find out as you try to pass the tests.

Good luck with your task!

Now you are all pumped up and ready to make an impact in the new company. You spend the afternoon with your mentor, who goes through the directory structure and the development procedure with you:

1. `jbod.h`: The interface of JBOD. You will use the constants defined here in your implementation.

2. `jbod.o`: The object file containing the JBOD driver.
3. `mdadm.h`: A header file that lists the functions you should implement.
4. `mdadm.c`: Your implementation of `mdadm` functions.
5. `tester.h`: Tester header file.
6. `tester.c`: Unit tests for the functions that you will implement. This file will compile into an executable, `tester`, which you will run to see if you pass the unit tests.
7. `util.h`: Utility functions used by JBOD implementation and the tester.
8. `util.c`: Implementation of utility functions.
9. `Makefile`: instructions for compiling and building `tester` used by the `make` utility.

You workflow will consist of (1) implementing functions by modifying `mdadm.c`, (2) typing `make` to build the `tester`, and (3) running `tester` to see if you pass the unit tests, and repeating these three steps until you pass all the tests. Although you only need to edit `mdadm.c` for successfully completing the assignment, you can modify any file you want if it helps you in some way. When testing your submission, however, we will use the original forms of all files except `mdadm.c` and `mdadm.h`.